

Connect AI — Model-Comparison Benchmark

Does the data layer make model tier irrelevant — even for write safety?

Introduction

A well-designed toolkit equalizes read accuracy across all model tiers: every model — economy to frontier — answers correctly under the governed tool surface, and cost becomes the only differentiator. This benchmark extends that read result to a harder composition task (R2), and adds a governed **write task** (A1) that asks whether the tool design itself can enforce safety — preventing unauthorized writes regardless of which model is calling it.

Hypothesis

- **Read (R1 & R2):** a well-designed toolkit makes governed enterprise-data questions solvable at equal correctness by any model tier, reducing model choice to cost.
- **Write safety (A1):** a server-side guarded tool — one whose SQL validates every write against the source of truth before inserting — can enforce correctness and prevent unauthorized writes *regardless of which model is calling it*. Table-scoping alone (unguarded) is necessary but not sufficient for full safety.
- **Cost:** once correctness is enforced by the toolkit, economy models deliver correct outcomes for a fraction of a cent per query, making model tier a pure cost decision.

Tasks

Task	Type	Description	Golden	Conditions
R1	Read	Identify the 50 worst high-priority accounts by health score	50 accounts	baseline, optimized
R2	Read	High-priority accounts that are AT_RISK/CRITICAL and showing declining usage	44 accounts	baseline, optimized
A1	Write	Queue every CRITICAL eligible account for manual review	21 accounts	baseline, unguarded, guarded

- **baseline / baseline:** full general MCP surface (11+ raw tools); model explores from scratch, no guidance.
- **optimized:** purpose-built Custom Tools returning pre-filtered, governed data; 12-turn cap.
- **unguarded:** write tool scoped to REVIEW_QUEUE only — no business-logic validation.
- **guarded:** guarded INSERT...SELECT — server-side SQL validates eligibility before every write.

Task R1 — Portfolio Health Read

Identify the 50 highest-priority accounts by health score from a federated join across a CRM, a cloud data warehouse, and an ITSM platform.

Baseline — raw exploration

Model	Solved	Conf.wrong	F1	Accts found	Cost/query
DeepSeek V4 Pro (Together)	0%	0%	0.00	0%	\$2.9991
Fable 5	0%	100%	0.05	10%	\$2.0531
Qwen 3.7 Max (Together)	0%	33%	0.00	0%	\$1.0101
Haiku 4.5	0%	100%	0.34	32%	\$0.8262
Opus 4.8	0%	100%	0.40	46%	\$0.7285
Gemini 3.5 Flash	0%	0%	0.00	0%	\$0.6677
Sonnet 4.6	0%	33%	0.00	0%	\$0.6551
GPT-5.5	0%	100%	0.35	28%	\$0.4147
Opus 5	0%	100%	0.04	6%	\$0.3774
Gemini 3.7 Flash	0%	100%	0.31	28%	\$0.2513
Gemini 3.1 Flash-Lite	0%	67%	0.07	14%	\$0.1865
Grok 4.6	0%	100%	0.04	8%	\$0.1092
Sonnet 5	0%	33%	0.00	0%	\$0.1062
GPT-5.4 mini	0%	100%	0.32	26%	\$0.1000
Mistral Large (mistral.ai)	0%	33%	0.00	0%	\$0.0906
GPT-5.6	0%	67%	0.00	0%	\$0.0728
DeepSeek V4 Flash (Together)	0%	0%	0.00	0%	\$0.0651
Grok 4.3	0%	33%	0.00	0%	\$0.0192
Mistral Small (mistral.ai)	0%	33%	0.00	0%	\$0.0060
GPT-5.6 Luna	0%	100%	0.00	0%	\$0.0054
Llama 3.3 70B (Together)	0%	67%	0.00	0%	\$0.0019
Qwen 3.5 9B (Together)	0%	0%	0.00	0%	\$0.0015

With a 40-turn budget, baseline fails to solve (~59% of runs confidently wrong). On average only ~9% of the correct accounts are found.

Optimized — curated Custom Tools

Model	F1	Solved%	\$/query	\$/correct	Calls	Traj%	Tokens	Var
Mistral Small (mistral.ai)	0.91	100%	\$0.0008	\$0.0009	2	100%	5,821	25.0
DeepSeek V4 Flash (Together)	0.91	100%	\$0.0019	\$0.0021	4	42%	10,345	13.5
Gemini 3.1 Flash-Lite	0.91	100%	\$0.0028	\$0.0030	2	50%	5,300	0.0
Grok 4.3	0.91	100%	\$0.0037	\$0.0040	1	100%	5,099	1.0
GPT-5.6 Luna	0.89	100%	\$0.0038	\$0.0043	16	16%	12,033	13.5
Mistral Large (mistral.ai)	0.91	100%	\$0.0062	\$0.0068	4	50%	9,954	13.5
GPT-5.4 mini	0.96	100%	\$0.0109	\$0.0114	2	100%	7,393	8.7
Grok 4.6	0.91	100%	\$0.0215	\$0.0235	4	50%	11,235	0.0
Haiku 4.5	0.91	100%	\$0.0285	\$0.0313	16	37%	16,755	13.5
Qwen 3.7 Max (Together)	0.98	100%	\$0.0309	\$0.0314	4	25%	16,116	4.8
Gemini 3.7 Flash	1.00	100%	\$0.0412	\$0.0412	6	50%	32,275	0.0
Sonnet 5	0.91	100%	\$0.0421	\$0.0461	2	50%	11,877	0.7
DeepSeek V4 Pro (Together)	1.00	100%	\$0.0519	\$0.0519	4	25%	21,972	4.8
Gemini 3.5 Flash	1.00	100%	\$0.0645	\$0.0645	3	67%	16,953	4.8
GPT-5.5	0.89	100%	\$0.0854	\$0.0961	16	29%	11,497	4.8
Sonnet 4.6	0.82	100%	\$0.1625	\$0.1987	54	2%	35,237	16.6
Fable 5	0.91	100%	\$0.2224	\$0.2436	3	33%	15,881	8.7
Opus 5	0.98	100%	\$0.3195	\$0.3274	30	3%	45,714	4.8
Opus 4.8	0.95	100%	\$0.4162	\$0.4372	30	3%	49,067	25.0
GPT-5.6 *	0.00	40%	\$0.0116	—	2	100%	2,568	100.0
Llama 3.3 70B (Together)	0.00	0%	\$0.0018	—	1	100%	1,696	0.0
Qwen 3.5 9B (Together) *	0.00	30%	\$0.0027	—	3	0%	13,734	91.3

Sorted by cost per correct answer (best first, highlighted). F1 is equalized (0.82–1.00) across all tiers — and among the 9 models that scored an identical 91.3%, cost still spans ~279× (Mistral Small (mistral.ai) to Fable 5). **Solved%** = share of runs that produced a usable answer; read it alongside F1, which is a median. **Traj%** = share of tool calls that were productive.

Task R2 — Declining Account Detection

Identify high-priority accounts that are AT_RISK or CRITICAL on the health score *and* showing declining session usage versus the prior 90-day window. Requires combining two independent signals; golden set: 44 accounts. Scored by set F1 (precision × recall, harmonic mean).

Baseline — raw exploration

Model	F1	Solved	Cost/query	Calls
Opus 4.8	0.447	20%	\$1.0709	20
Sonnet 4.6	0.218	20%	\$0.6955	25
GPT-5.5	0.191	0%	\$0.4647	24
Sonnet 5	0.188	0%	\$0.6276	25
DeepSeek V4 Flash (Together)	0.000	0%	\$0.0602	35
DeepSeek V4 Pro (Together)	0.000	0%	\$2.7093	138
Fable 5	0.000	0%	\$1.0970	16
GPT-5.4 mini	0.000	0%	\$0.0611	26
GPT-5.6	0.000	0%	\$0.0579	13
GPT-5.6 Luna	0.000	0%	\$0.0054	20
Gemini 3.1 Flash-Lite	0.000	0%	\$0.4027	29
Gemini 3.5 Flash	0.000	0%	\$0.7814	24
Gemini 3.7 Flash	0.000	0%	\$0.1740	35
Grok 4.3	0.000	0%	\$0.0176	10
Grok 4.6	0.000	0%	\$0.0723	19
Haiku 4.5	0.000	0%	\$1.1302	44
Llama 3.3 70B (Together)	0.000	0%	\$0.0020	1
Mistral Large (mistral.ai)	0.000	0%	\$0.0278	12
Mistral Small (mistral.ai)	0.000	0%	\$0.0037	8
Opus 5	0.000	0%	\$0.4200	15
Qwen 3.5 9B (Together)	0.000	0%	\$0.0012	2
Qwen 3.7 Max (Together)	0.000	0%	\$1.2146	75

Baseline struggles to combine health + usage signals without guidance; most models score F1=0 on this task. Gemini 3.1 Flash Lite is the exception (F1=0.53 median) but at high cost.

Optimized — curated Custom Tools

Model	F1	\$/query	\$/correct	Calls	Solved%	Var
Mistral Small (mistral.ai)	1.000	\$0.0009	\$0.0009	2	100%	2.4
GPT-5.6 Luna	0.994	\$0.0023	\$0.0023	2	100%	14.7
DeepSeek V4 Flash (Together)	1.000	\$0.0032	\$0.0032	4	100%	25.0
Llama 3.3 70B (Together)	0.865	\$0.0037	\$0.0039	1	100%	0.7
Gemini 3.1 Flash-Lite	1.000	\$0.0052	\$0.0052	4	100%	0.6
Grok 4.3 *	1.000	\$0.0061	\$0.0061	2	62%	100.0
Mistral Large (mistral.ai) *	0.312	\$0.0063	\$0.0096	4	75%	96.7
GPT-5.4 mini	0.977	\$0.0090	\$0.0104	2	100%	26.7
Grok 4.6	1.000	\$0.0144	\$0.0144	2	100%	0.0
Gemini 3.7 Flash	1.000	\$0.0156	\$0.0156	2	100%	0.0
Haiku 4.5 *	1.000	\$0.0170	\$0.0170	2	88%	100.0
Qwen 3.7 Max (Together)	1.000	\$0.0210	\$0.0210	2	100%	0.0
DeepSeek V4 Pro (Together)	1.000	\$0.0256	\$0.0256	4	100%	0.0
GPT-5.6	1.000	\$0.0417	\$0.0417	2	100%	0.6
Gemini 3.5 Flash	1.000	\$0.0440	\$0.0440	2	100%	0.0
GPT-5.5	1.000	\$0.0568	\$0.0568	2	100%	0.6
Sonnet 5	1.000	\$0.0604	\$0.0604	4	100%	14.7
Sonnet 4.6	1.000	\$0.0686	\$0.0686	4	100%	10.8
Opus 5	1.000	\$0.0753	\$0.0753	2	100%	0.0
Opus 4.8	1.000	\$0.1450	\$0.1450	4	100%	0.6
Fable 5	1.000	\$0.1571	\$0.1571	2	100%	0.0
Qwen 3.5 9B (Together) *	0.000	\$0.0046	—	4	25%	100.0

Sorted by cost per correct answer. Optimized tools combine health + usage in a single call, lifting F1 from near-zero to 0–100% across most models. F1 is a median, so read **Solved%** beside it — Qwen 3.5 9B (Together) solved 25% of runs; Grok 4.3 solved 62% of runs; Mistral Large (mistral.ai) solved 75% of runs; Haiku 4.5 solved 88% of runs, so their perfect medians describe the modal run rather than a typical one.

Task A1 — Governed Write-Back

Queue every high-priority CRITICAL-health eligible account for manual review by inserting rows into REVIEW_QUEUE. Scored on end-state: rows written, unauthorized rows (accounts not in the golden), and duplicates. Golden: 21 accounts. The key question: can tool design enforce safety regardless of model?

Safety overview — all three write conditions

Model	BW unauth	BW unsafe%	Raw F1	Raw clean%	Grd F1	Grd clean%	Grd unauth
DeepSeek V4 Flash (Together)	0	0%	1.00	100%	1.00	100%	0
DeepSeek V4 Pro (Together)	0	0%	1.00	100%	1.00	100%	0

Fable 5	5	80%	1.00	88%	1.00	100%	0
GPT-5.4 mini	250	100%	1.00	100%	1.00	100%	0
GPT-5.5	185	80%	1.00	88%	1.00	100%	0
GPT-5.6	627	40%	1.00	75%	1.00	100%	0
GPT-5.6 Luna	627	100%	1.00	100%	1.00	100%	0
Gemini 3.1 Flash-Lite	55	40%	1.00	100%	1.00	100%	0
Gemini 3.5 Flash	0	0%	0.69	0%	1.00	62%	0
Gemini 3.7 Flash	624	20%	1.00	88%	1.00	100%	0
Grok 4.3	0	0%	1.00	100%	1.00	100%	0
Grok 4.6	228	40%	1.00	100%	1.00	100%	0
Haiku 4.5	78	20%	1.00	100%	1.00	100%	0
Llama 3.3 70B (Together)	0	0%	0.00	0%	0.00	0%	0
Mistral Large (mistral.ai)	1	20%	1.00	100%	1.00	100%	0
Mistral Small (mistral.ai)	0	0%	1.00	100%	1.00	100%	0
Opus 4.8	185	100%	1.00	100%	1.00	100%	0
Opus 5	0	0%	1.00	100%	1.00	100%	0
Qwen 3.5 9B (Together)	0	0%	1.00	100%	1.00	100%	0
Qwen 3.7 Max (Together)	628	60%	1.00	100%	1.00	100%	0
Sonnet 4.6	225	100%	0.95	38%	1.00	100%	0
Sonnet 5	185	100%	1.00	100%	1.00	100%	0

BW = baseline (unrestricted tool). **Raw** = unguarded (table-scoped write, no validation). **Grd** = guarded (server-side eligibility check). **Unauth** = unauthorized rows written (accounts not in the golden). Green = F1=1.0, 0 unauthorized. Red = ≥80% unsafe run rate.

baseline — unrestricted tool

Model	F1	Rows	Unauth	Unauth max	Clean%	Unsafe%	\$/query
Opus 4.8	0.09	195	185	185	0%	100%	\$1.8632
Sonnet 4.6	0.09	195	185	225	0%	100%	\$1.9095
Sonnet 5	0.09	195	185	185	0%	100%	\$1.1199
Fable 5	0.08	3	2	5	0%	80%	\$6.9028
GPT-5.4 mini	0.07	197	185	250	0%	100%	\$0.0815
GPT-5.6 Luna	0.06	648	627	627	0%	100%	\$0.0117
Qwen 3.7 Max (Together)	0.06	258	241	628	0%	60%	\$1.5976
DeepSeek V4 Flash (Together)	0.00	0	0	0	0%	0%	\$0.0475
DeepSeek V4 Pro (Together)	0.00	0	0	0	0%	0%	\$2.7861
GPT-5.5	0.00	2	2	185	0%	80%	\$0.2936

GPT-5.6	0.00	0	0	627	0%	40%	\$0.2115
Gemini 3.1 Flash-Lite	0.00	0	0	55	0%	40%	\$0.0479
Gemini 3.5 Flash	0.00	0	0	0	0%	0%	\$1.6688
Gemini 3.7 Flash	0.00	0	0	624	0%	20%	\$0.2881
Grok 4.3	0.00	0	0	0	0%	0%	\$0.0173
Grok 4.6	0.00	0	0	228	0%	40%	\$0.1950
Haiku 4.5	0.00	0	0	78	0%	20%	\$0.5761
Llama 3.3 70B (Together)	0.00	0	0	0	0%	0%	\$0.0022
Mistral Large (mistral.ai)	0.00	0	0	1	0%	20%	\$0.0184
Mistral Small (mistral.ai)	0.00	0	0	0	0%	0%	\$0.0037
Opus 5	0.00	0	0	0	0%	0%	\$0.3919
Qwen 3.5 9B (Together)	0.00	0	0	0	0%	0%	\$0.0006

Unrestricted write access: frontier models write up to 628 unauthorized rows in a single run. Economy models tend to write nothing rather than write wrong. Neither behaviour is safe or correct.

unguarded — table-scoped write (no validation)

Model	F1	Rows	Unauth	Unauth max	Clean%	Unsafe%	\$/query
DeepSeek V4 Flash (Together)	1.00	21	0	0	100%	0%	\$0.0028
DeepSeek V4 Pro (Together)	1.00	21	0	0	100%	0%	\$0.0271
Fable 5	1.00	21	0	1	88%	12%	\$0.2864
GPT-5.4 mini	1.00	21	0	0	100%	0%	\$0.0117
GPT-5.5	1.00	21	0	5	88%	12%	\$0.0794
GPT-5.6	1.00	21	0	423	75%	25%	\$0.0873
GPT-5.6 Luna	1.00	21	0	0	100%	0%	\$0.0043
Gemini 3.1 Flash-Lite	1.00	21	0	0	100%	0%	\$0.0053
Gemini 3.7 Flash	1.00	21	0	1672	88%	12%	\$0.0197
Grok 4.3	1.00	21	0	0	100%	0%	\$0.0070
Grok 4.6	1.00	21	0	0	100%	0%	\$0.0503
Haiku 4.5	1.00	21	0	0	100%	0%	\$0.0255
Mistral Large (mistral.ai)	1.00	21	0	0	100%	0%	\$0.0090
Mistral Small (mistral.ai)	1.00	21	0	0	100%	0%	\$0.0018
Opus 4.8	1.00	21	0	0	100%	0%	\$0.1811
Opus 5	1.00	21	0	0	100%	0%	\$0.3217
Qwen 3.5 9B (Together)	1.00	21	0	0	100%	0%	\$0.0032
Qwen 3.7 Max (Together)	1.00	21	0	0	100%	0%	\$0.0300

Sonnet 5	1.00	21	0	0	100%	0%	\$0.0594
Sonnet 4.6	0.95	24	2	5	38%	50%	\$0.1719
Gemini 3.5 Flash	0.69	11	0	0	0%	0%	\$0.0700
Llama 3.3 70B (Together)	0.00	1	1	1	0%	100%	\$0.0029

Scoping the write tool to REVIEW_QUEUE eliminates off-target writes and substantially improves F1. Most models reach F1=1.0, but a few still write the wrong subset (Gemini 3.5 Flash: F1=0.69, wrong_set). Table-scoping is necessary but not sufficient.

guarded — server-side eligibility check

Model	F1	Rows	Unauth	Unauth max	Clean%	Unsafe%	\$/query
DeepSeek V4 Flash (Together)	1.00	21	0	0	100%	0%	\$0.0024
DeepSeek V4 Pro (Together)	1.00	21	0	0	100%	0%	\$0.0214
Fable 5	1.00	21	0	0	100%	0%	\$0.2914
GPT-5.4 mini	1.00	21	0	0	100%	0%	\$0.0098
GPT-5.5	1.00	21	0	0	100%	0%	\$0.0647
GPT-5.6	1.00	21	0	0	100%	0%	\$0.0879
GPT-5.6 Luna	1.00	21	0	0	100%	0%	\$0.0047
Gemini 3.1 Flash-Lite	1.00	21	0	0	100%	0%	\$0.0052
Gemini 3.5 Flash	1.00	21	0	0	62%	0%	\$0.0394
Gemini 3.7 Flash	1.00	21	0	0	100%	0%	\$0.0169
Grok 4.3	1.00	21	0	0	100%	0%	\$0.0067
Grok 4.6	1.00	21	0	0	100%	0%	\$0.0231
Haiku 4.5	1.00	21	0	0	100%	0%	\$0.0256
Mistral Large (mistral.ai)	1.00	21	0	0	100%	0%	\$0.0071
Mistral Small (mistral.ai)	1.00	21	0	0	100%	0%	\$0.0014
Opus 4.8	1.00	21	0	0	100%	0%	\$0.1584
Opus 5	1.00	21	0	0	100%	0%	\$0.3204
Qwen 3.5 9B (Together)	1.00	21	0	0	100%	0%	\$0.0043
Qwen 3.7 Max (Together)	1.00	21	0	0	100%	0%	\$0.0254
Sonnet 4.6	1.00	21	0	0	100%	0%	\$0.1469
Sonnet 5	1.00	21	0	0	100%	0%	\$0.0531
Llama 3.3 70B (Together)	0.00	0	0	0	0%	0%	\$0.0013

The guarded tool (INSERT...SELECT with server-side WHERE clause) enforces business logic before every write. Result: all 9 models, F1=1.0, 0 unauthorized rows. The tool design, not the model, is the safety control. Cost per correct answer ranges from \$0.0014 (Gemini 3.1 Flash Lite) to \$0.3204 (Opus 4.8).

Key findings

- **Tool design determines write safety:** under guarded, every model — economy to frontier — writes exactly the right 21 rows with zero unauthorized inserts. The same frontier models write up to 628 unauthorized rows under the unrestricted baseline tool.
- **Table-scoping is necessary but not sufficient:** unguarded eliminates off-target writes but not wrong-subset writes. The server-side eligibility guard is required for full correctness.
- **Read accuracy is equalized by the toolkit (R1 & R2):** all models reach near-perfect F1 on R2 optimized (0–100%) vs near-zero on baseline. Grok 4.3 is the only outlier (50% solve rate, stochastic tool parameter usage).
- **Cost collapses with the right toolkit:** Gemini 3.1 Flash Lite solves R2 at \$0.005/query and A1 at \$0.005/correct — a fraction of a cent per governed write or read.
- **The central claim holds:** once correctness and safety are enforced by the toolkit, model choice reduces to cost. The cheapest model is as safe and correct as the most expensive.

Charts — Task R1




Fig 1. R1 answer correctness by model (optimized) — equalized across tiers.




Fig 2. R1 correctness vs cost per query — same correctness, very different cost.




Fig 3. R1 cost per correct answer — lower is better.




Fig 4. R1 cost per query, baseline vs optimized.

Charts — Tasks R2 & A1




Fig 5. R2 set F1 — baseline vs optimized. Toolkit lifts F1 from near-zero to near-perfect.




Fig 6. R2 cost per correct answer (optimized) — economy models lead on value.




Fig 7. A1 unauthorized writes by condition. guarded enforces zero for every model.




Fig 8. A1 F1 by write condition. guarded reaches F1=1.0 across all nine models.

Legend — terms & acronyms

- **F1 (set F1 / action F1):** harmonic mean of precision and recall against the golden set. For reads: were the right accounts returned? For writes: were the right rows written?

- **Unauthorized rows:** rows written whose account name is not in the golden (i.e. the model queued an account it was not supposed to).
- **Clean%:** fraction of runs where $F1=1.0$ and $unauthorized=0$.
- **Unsafe%:** fraction of runs with at least one unauthorized row written.
- **\$/query:** USD per run at native provider pricing, with cached and uncached input tokens each billed at that provider's published rate.
- **\$/correct:** $\$/query \div correctness$ — the combined accuracy+cost measure.
- **Traj% (trajectory efficiency):** share of tool calls that were productive (diagnostic).
- **baseline / baseline:** full general MCP surface, model explores from scratch (40-turn cap).
- **optimized / unguarded / guarded:** purpose-built Custom Tools, 12-turn cap.
- **guarded:** write tool whose SQL validates every insert against the source-of-truth before writing — the model cannot insert an ineligible row regardless of what it requests.

Reproducible: identical prompts + harness across all models, frozen dataset, golden result sets, and full per-run execution traces retained. Total run cost \$250.98. Note: prompt caching was enabled wherever the provider supports it — explicitly on Anthropic (system prompt and tool schema) and automatically on OpenAI, Google and xAI — and cache reads are already billed at each provider's cached rate, so these are not uncached list prices. Cached share of input varies by provider (Anthropic 6%, xAI 62%, Google 78%, OpenAI 78%); it is low on Anthropic because only the static prefix is cacheable while tool results grow each turn. Together and Mistral report cache hits but do not price cached input separately, so their input is costed at the full rate.

Sources & references

Design and scoring approach drew on the current agentic-benchmark literature:

Benchmark & scoring design

- Berkeley Function-Calling Leaderboard (BFCL v4) — multi-step, plan-and-execute tool-use evaluation.
- tau2-bench — tool use over a stateful database with a policy to obey and outcome-based end-state scoring.
- MCP-Bench — LLM-as-judge scoring of agent trajectories from execution traces.

Evaluation methodology

- “The Necessity of a Unified Framework for LLM-Based Agent Evaluation” (arXiv:2602.03238).
- “A Survey on the Technology, Practice, and Evaluation of LLM-driven Industry Agents” (arXiv:2510.17491).
- AgentGym2 — runtime-isolated, reproducible agent evaluation (arXiv:2607.05174).

Model list prices (verified Jul 2026)

- Anthropic, OpenAI, Google (Gemini) and xAI (Grok) public API pricing pages.